

qemu kvm study

<https://lwn.net/Articles/810033/>

1. EPT guest paging structures → Host physical address guest-physical address: Only 47:0

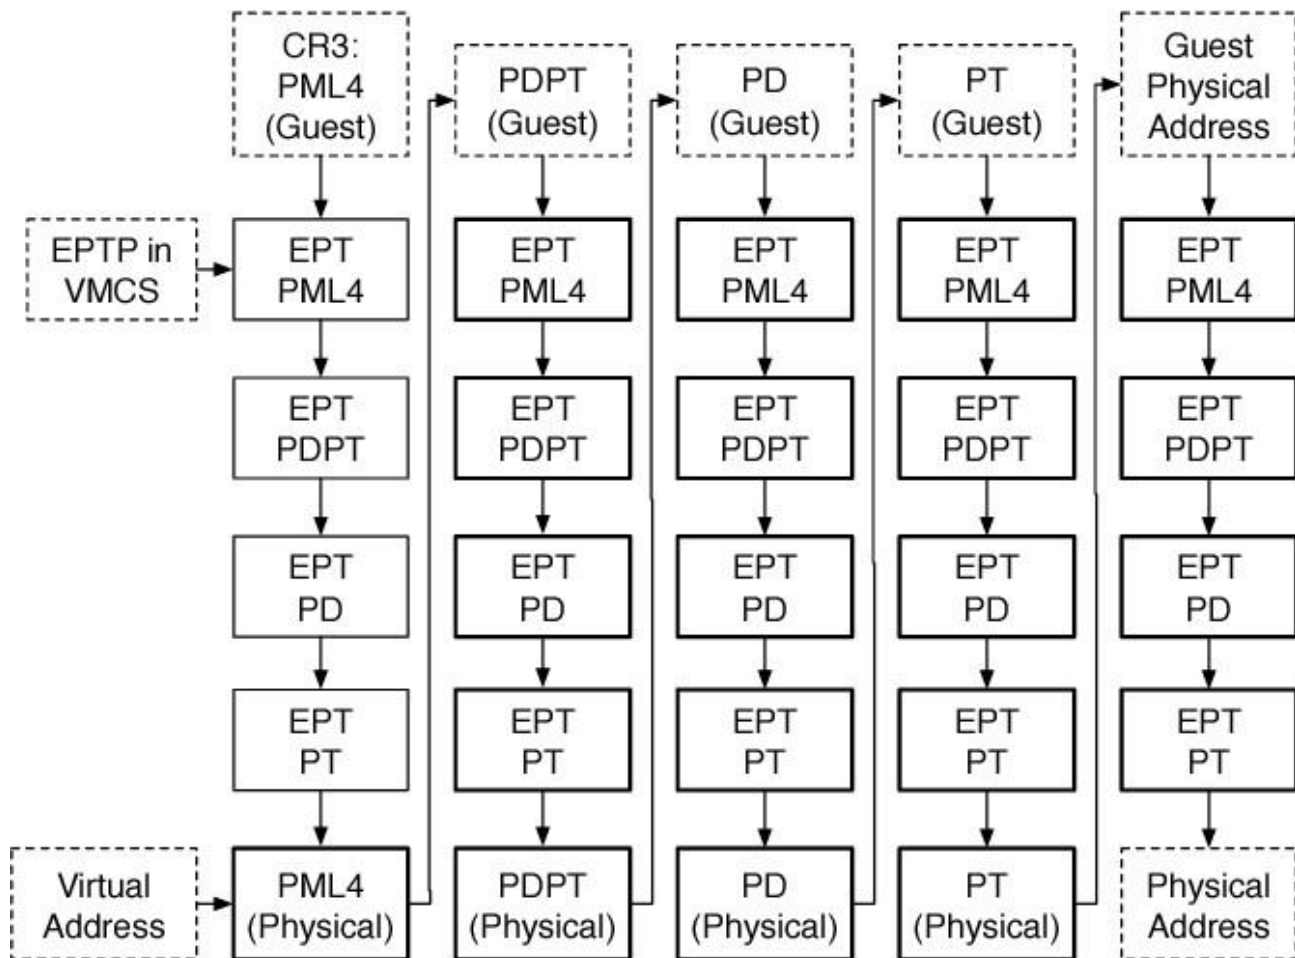
will be used for 64 bit address; high bit should be zero,

4 级页表:

47:0	<u>47: 39</u>	[38:30]	[29:21]	[20:12]	low 12bit
最大 256T	512G/entry	1G/entry	2M/entry	4k/entry	4K
EPTP	EPT PML4 Entry[512]	EPT Page-DirectoryPointer[512]	EPT page table[512]	EPT pte[512]	one host page

63	62	61	60	59	58	57	56	55	54	53	52	51	50		M ¹	M-1			31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0							
Reserved																Address of EPT PML4 table																Rsvd.		SSS		A/D		EPT PML-1		EPT PS MT		EPTP															
Ignored												Rsvd.				Address of EPT page-directory-pointer table																lg n.		X4		lg n.		A		Reserved		X5		W		R		PML4 present									
Ignored																Ignored																								PML4 not present																	
Ignored																Rsvd.				Physical address of 1GB page				Reserved												lg n.		X4		D		A		1		IPAT		EPT MT		X		W		R		PDPT 1GB page	
Ignored												Rsvd.				Address of EPT page directory																lg n.		X4		lg n.		A		0		Rsvd.		X		W		R		PDPT page directory							
Ignored																Ignored																								PDPT not present																	
Ignored																Rsvd.				Physical address of 2MB page								Reserved								lg n.		X4		D		A		1		IPAT		EPT MT		X		W		R		PDE 2MB page	
Ignored												Rsvd.				Address of EPT page table																lg n.		X4		lg n.		A		0		Rsvd.		X		W		R		PDE page table							
Ignored																Ignored																								PDE not present																	
Ignored																Rsvd.				Physical address of 4KB page																lg n.		X4		D		A		lg n.		IPAT		EPT MT		X		W		R		PTE 4KB page	
Ignored																Ignored																								PTE not present																	

physical address, and CR3, and EPTP in VMCS and all configured:



TLB (bigger) and Paging-Structure Caches

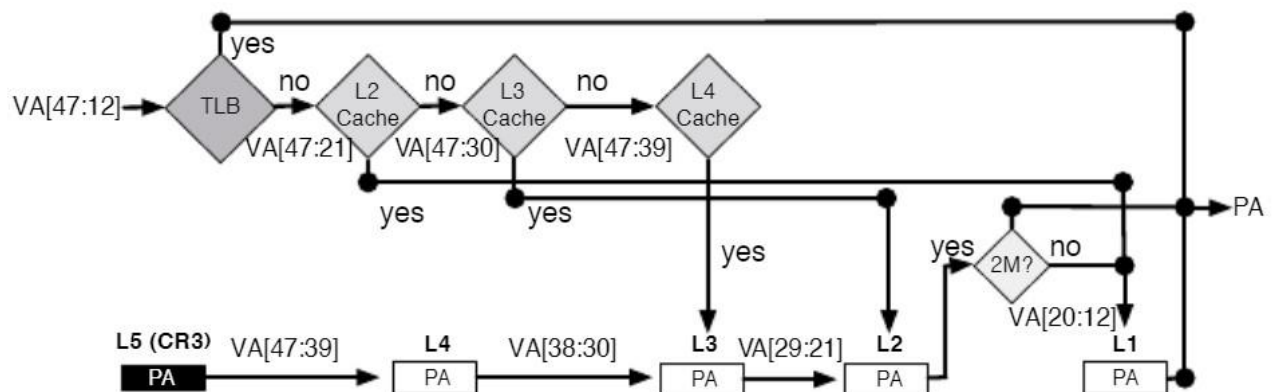


Figure 1. The x86 native page walk.

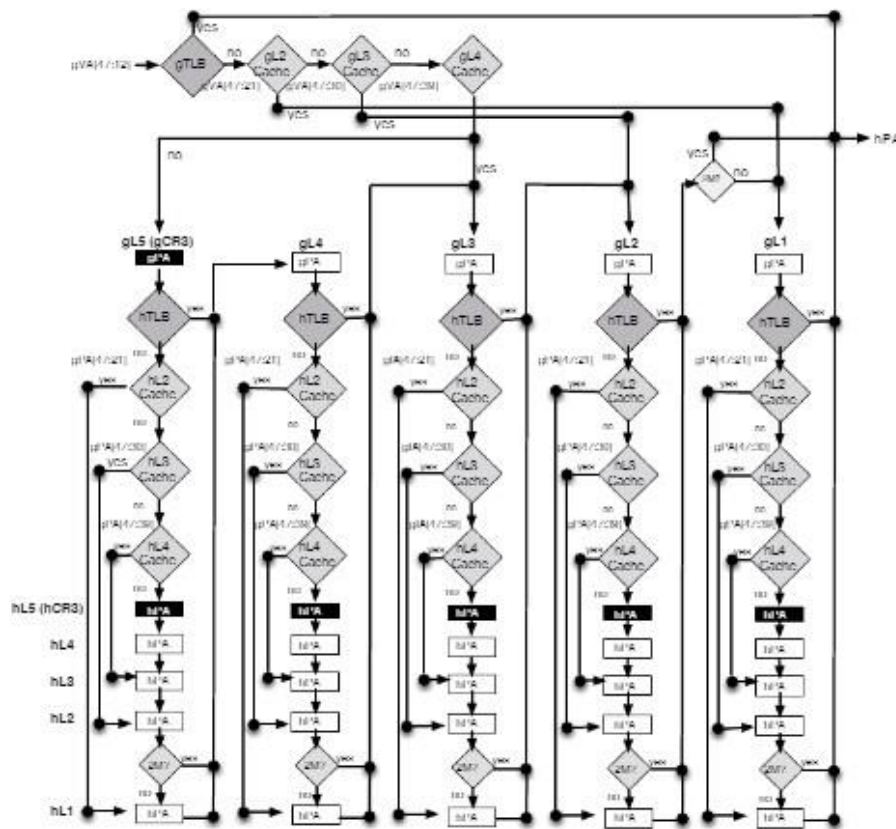


Figure 2. The x86 2-D (virtual) page walk.

EPT induced VM exit:

EPT misconfigurations	EPT paging-structure entry that contains an unsupported value	Bit 0 is clear bit 1 is set	
EPT violations	EPT pagingstructure entries disallow an access using the guest-physical address	an EPT paging-structure entry that is not present	
pagemodification log-full event	the logical processor determines a need to create a pagemodification log entry and the current log is full		

Accessed and Dirty Flags for EPT

Whenever the processor uses an EPT paging-structure entry as part of guest-physical-address translation, it sets the accessed flag in that entry		
Whenever there is a write to a guest-physical address, the processor sets the dirty flag		
These flags are “sticky,” meaning that, once set, the processor does not clear them; only software can clear them.		

Page-Modification Logging Memory type:

EPT 页表的 memory type:

If CR0.CD = 0, ept 页表的 memory type 取决于 extended-page-table pointer (EPTP) 2:0 , 0--uncacheable type (UC),
6write-back type
(WB)

If CR0.CD = 0, ept 页表的 memory type, uncacheable (UC);

Memory type for Translated Guest-Physical Addresses 三个因素:

- 1.CR0 CD;
- 2.last EPT paging-structure entry;
3. PAT memory type;

无标签